

# Report: Recommender Systems — Part 2

Building on Part 1, Part 2 addresses the core limitations identified there: objective mismatch (MF models optimize MSE, not ranking quality), cold-start (CF models need interaction history), and narrow coverage (accurate models collapse into a small subset of popular items). We implement ranking heuristics, pairwise BPR, hybrid architectures, neural models, and bandit-based online evaluation.

All models share the same evaluation protocol from Part 1: per-user temporal split 80/10/10, relevance threshold  $\geq 4.0$ , NDCG@K as primary metric,  $K \in \{5, 10, 20\}$ .

## 1. Ranking Heuristics

Three non-learned baselines that require no training.

### 1.1 Popularity Ranker

Counts all interactions per item (`train_df.groupby('item_id').size()`). Every rating counts — a 1-star and a 5-star rating contribute equally. The item's score is simply "how many people rated this movie." Minimum threshold: `min_ratings=5`.

### 1.2 Recency Ranker

Each interaction gets a time-decay weight:  $\exp(-\ln 2 \cdot \text{days\_ago} / \text{half\_life})$  with `half_life=90` days. A rating from yesterday weighs  $\approx 1.0$ ; a rating from 180 days ago weighs  $\approx 0.25$ . The actual rating value (1–5) is ignored — only the timestamp matters. The item's score is the sum of all its time-decayed weights. On MovieLens with only a 3-year temporal span, recency added little over raw popularity.

### 1.3 Personalized PageRank (PPR)

Builds a bipartite graph: nodes = all users + all items, edges = "user u liked item i" (rating  $\geq 4.0$ ). Row-normalizes the adjacency matrix to get a transition matrix (each node distributes 1/degree to each neighbor). Then runs power iteration for 20 steps with teleport probability  $\alpha=0.15$ :

```
scores = (1 - alpha) * M^T * scores + alpha * teleport
```

where `teleport` is 1.0 at the target user and 0.0 elsewhere. The result: every item reachable through any path from the user gets a score — even items the user never interacted with, reached through multi-hop chains (user→movie→user→movie).

## Heuristic Results

Model	NDCG@5	NDCG@10	NDCG@20	Precision@10	Recall@10	Coverage	Pop. Bias
PPR	<b>0.0475</b>	<b>0.0544</b>	<b>0.0667</b>	<b>0.0428</b>	<b>0.0497</b>	5.8%	2053
Popularity	0.0437	0.0490	0.0615	0.0393	0.0433	5.1%	2086
Recency	0.0380	0.0438	0.0547	0.0353	0.0410	4.4%	1797

PPR is the best heuristic — it captures indirect user-item connections through the graph (e.g., "users who liked the same movies as you also liked this"). But all heuristics have extremely low coverage (<6%) because they collapse into the popular, well-connected core.

## 2. BPR-MF: Pairwise Learning-to-Rank

### The core idea

BPR uses the same latent factor architecture as FunkSVD (user vectors P, item vectors Q, biases) but replaces the MSE loss with a pairwise ranking loss: for each training step, sample a user, one of their liked items (positive), and a random item they haven't interacted with (negative), then push  
 $\text{score}(\text{user}, \text{positive}) > \text{score}(\text{user}, \text{negative})$ .

The loss per sample:  $-\log \sigma(x_{ui} - x_{uj})$  where  $\sigma$  is the sigmoid function.

### Implementation details

- **Architecture:** n\_factors=64, user\_bias + item\_bias
- **Initialization:** user\_factors and item\_factors from  $N(0, 0.01)$ , biases at 0
- **Training:** lr=0.01, reg=0.001, 40 epochs
- **Negative sampling:** n\_samples = len(positive\_interactions)  $\times$  5 per epoch. Each sample: pick random user → pick one of their positives → sample one random negative.
- **Acceleration:** Full training loop compiled via Numba `@njit(fastmath=True)`
- **Scoring:** `score_all_items(u) = item_bias + item_factors @ user_factors[u]` — a single matrix-vector multiply producing 3,706 scores

## BPR vs. Baselines

Model	NDCG@5	NDCG@10	NDCG@20	Precision@10	Recall@10	Coverage	Pop. Bias
<b>BPR-MF</b>	<b>0.0567</b>	<b>0.0664</b>	<b>0.0829</b>	<b>0.0466</b>	<b>0.0704</b>	<b>45.9%</b>	1095
PPR	0.0475	0.0544	0.0667	0.0428	0.0497	5.8%	2053
Popularity	0.0437	0.0490	0.0615	0.0393	0.0433	5.1%	2086

BPR outperforms all heuristic baselines and all Part 1 models (best Part 1: Item-Item CF at NDCG@10 = 0.0625) while dramatically improving coverage. Compared to FunkSVD (same architecture, MSE loss): NDCG@10 jumps from 0.0265 → 0.0664 — a **+150% improvement** from switching the loss function alone. This is the single most impactful design choice in the project.

## Negative Sampling Sensitivity

Multiplier	Samples/epoch	NDCG@10	Recall@10	Coverage
1x	469,223	0.0502	0.0449	2.9%
<b>5x</b>	<b>2,346,115</b>	<b>0.0661</b>	<b>0.0700</b>	<b>34.2%</b>
10x	4,692,230	0.0651	0.0698	42.4%
20x	9,384,460	0.0646	0.0697	46.9%

At 1x the model sees too little contrast — each positive gets roughly one gradient update per epoch. At 5x each positive is sampled ~5 times (paired with different random negatives), giving rich contrast. Beyond 5x accuracy slightly degrades while coverage keeps increasing — the model learns to distribute scores across more items but loses some precision at the top.

## Head / Mid / Tail Analysis

Items split by popularity quantiles: head = top 10% ( $\geq q90$ ), mid = 50th–90th percentile, tail = bottom 50%.

Model	Head NDCG@10	Mid NDCG@10	Tail NDCG@10
<b>BPR-MF</b>	<b>0.0884</b>	<b>0.0238</b>	<b>0.0012</b>
PPR	0.0737	0.0	0.0
Popularity	0.0662	0.0	0.0

Popularity and PPR recommend exclusively head items — literally 0.0 on mid and tail. BPR is the only model that can surface mid-tier and tail items at all. The tail number (0.0012) is tiny but it is nonzero: BPR occasionally

recommends a niche movie that happens to be in the user's test set. This happens because negative sampling forces the model to learn embeddings for tail items — their positions in latent space are shaped by the gradient updates even though few users interacted with them directly.

## 3. Hybrid Recommender

### Motivation

BPR struggles with cold-start (users/items with few interactions). Content-Based can score any item with metadata but misses collaborative signal. The hybrid combines both.

### Two strategies implemented

**Weighted Blending:** Score all 3,706 items with both models. Min-max normalize each score array to [0, 1] (necessary because BPR scores are unbounded while CB scores are bounded). Compute  
 $\text{final} = \alpha \cdot \text{BPR\_norm} + (1-\alpha) \cdot \text{CB\_norm}$ . Return top-K.

**Candidate Generation + Reranking:** BPR retrieves top-100 candidates. For each candidate:  
 $\text{combined} = 0.6 \cdot \text{bpr\_score} + 0.4 \cdot \text{cb\_score}$  (raw scores, not normalized). Return top-K from the reranked list. This is more computationally efficient — CB only scores 100 items instead of 3,706.

### Alpha Tuning (Weighted Blending, cold users only)

$\alpha$ (BPR weight)	NDCG@10	Recall@10
0.3	0.0642	0.0750
0.5	0.0720	0.0840
0.7	0.0785	0.0914
<b>0.8</b>	<b>0.0811</b>	<b>0.0938</b>
0.9	0.0805	0.0922

$\alpha=0.8$  is optimal — 80% BPR, 20% CB. Below 0.8 the content signal dilutes the collaborative signal that BPR already provides. Above 0.8 there's not enough content fallback for cold items.

### Hybrid Results (Full Test Set)

Model	NDCG@10	Precision@10	Recall@10	Coverage	Pop. Bias
Blend ( $\alpha=0.8$ )	<b>0.0657</b>	<b>0.0452</b>	<b>0.0708</b>	<b>48.1%</b>	1058

Model	NDCG@10	Precision@10	Recall@10	Coverage	Pop. Bias
CandGen+Rerank	0.0655	0.0455	0.0701	47.2%	1067
BPR-MF (pure)	0.0649	0.0455	0.0689	45.8%	1094
Content-Based	0.0359	0.0224	0.0382	55.0%	734

Blend and CandGen+Rerank are nearly identical (0.002 difference = noise). The CandGen+Rerank architecture's value is not empirical on this dataset — it's architectural: scalability (CB scores 100 items instead of 3,706) and modularity (swappable retrieval and reranking stages).

## Cold vs. Warm User Segmentation

Model	Cold NDCG@10 ( $\leq 30$ ratings)	Warm NDCG@10 ( $\geq 100$ ratings)
Blend ( $\alpha=0.8$ )	<b>0.0806</b>	0.0668
CandGen+Rerank	0.0769	0.0693
BPR-MF	0.0730	<b>0.0707</b>
Content-Based	0.0478	0.0312

For cold users the Hybrid boosts NDCG by +10% over pure BPR. Content features fill in when interaction history is sparse. For warm users pure BPR is actually better — when the collaborative signal is strong, injecting generic content features adds noise.

## 4. Deep Learning Models

### 4.1 NeuMF (Neural Collaborative Filtering)

Combines a GMF path (element-wise product of user/item embeddings) and an MLP path (concat, then feed-forward). Four embedding tables: gmf\_user, gmf\_item, mlp\_user, mlp\_item — all emb\_dim=32, initialized from  $N(0, 0.01)$ . MLP dims: [64, 32] with ReLU and Dropout(0.1). Linear layers Xavier-initialized. Output:  $\text{concat}(\text{GMF } 32\text{-dim}, \text{ MLP } 32\text{-dim}) \rightarrow \text{Linear}(64, 1)$ .

Training: BPR pairwise loss (same as BPR-MF for fair comparison), Adam optimizer, lr=0.001, weight\_decay=1e-5, batch\_size=2048, 30 epochs.

**Inference limitation:** NeuMF mixes user and item data inside the network — to score one user against all items requires 3,706 separate forward passes. This makes it impractical for production.

## 4.2 Two-Tower

Separate user and item encoder towers:

- **User tower:** Embedding(32) → Linear(32→64) → ReLU → Dropout(0.1) → Linear(64→64) → 64-dim output
- **Item tower:** [Embedding(32) ; TF-IDF features(~30-dim)] → concat → Linear(~62→64) → ReLU → Dropout(0.1) → Linear(64→64) → 64-dim output
- **Score:** dot product of the two 64-dim vectors

Same BPR pairwise loss, same optimizer settings as NeuMF.

**Key advantage:** The item tower takes content features (TF-IDF genre + decade) alongside the learned embedding. At inference, all 3,706 item representations are precomputed through the item tower once and cached. Per-user inference reduces to one user tower forward pass + a matrix-vector dot product — compatible with ANN search (FAISS) for fast serving.

## Deep Learning Results

Model	NDCG@5	NDCG@10	NDCG@20	Precision@10	Recall@10	Coverage	Pop. Bias
BPR-MF (baseline)	0.0571	0.0657	0.0828	0.0457	0.0692	45.4%	1099
Two-Tower	0.0553	0.0648	0.0800	0.0477	0.0666	54.1%	1199
NeuMF	0.0529	0.0619	0.0765	0.0447	0.0622	56.8%	1227

Both neural models slightly underperform the linear BPR-MF. Neural networks need large amounts of data to learn well: on MovieLens 1M with purely structural (ID-based) interactions, the extra MLP parameters learn noise rather than useful patterns. A well-regularized linear dot product generalizes better on a dataset of this size.

Two-Tower beats NeuMF because its item tower incorporates TF-IDF content features — a richer item representation. Both neural models achieve higher coverage (54–57%) than BPR-MF (45.4%), suggesting the MLP spreads scores across a wider range of items.

# 5. Online Evaluation with Multi-Armed Bandits

## Setup

Offline metrics evaluate models on frozen historical data — they cannot capture exploration value or feedback loops. We simulated online evaluation using multi-armed bandits.

**Arms:** 4 static recommendation policies — Popularity, BPR-MF, Hybrid, Random. Each arm is a function:  
`user_id → top-K recommendations`.

**Reward:** Precision@K —  $\text{hits} / K$  where a hit is a recommended item that appears in the user's test set with rating  $\geq 4.0$ . This is a continuous value in  $\{0, 0.1, 0.2, \dots, 1.0\}$ , not strictly binary.

**Bandit strategies:**  $\epsilon$ -greedy ( $\epsilon=0.1, \epsilon=0.3$ ), UCB1, Thompson Sampling. For Thompson Sampling, the internal update is binarized (`reward > 0 → success`), but the reported metrics use the raw precision values.

**Simulation:** 6,040 users arrive sequentially in random order. At each step, the bandit picks an arm, the arm generates top-10 recommendations, the reward is computed from the test set, and the bandit updates its beliefs.

## Bandit Results

Strategy	Avg Reward	Total Reward
Static: Hybrid	0.0487	293.9
Static: BPR-MF	0.0470	283.7
Static: Popularity	0.0393	237.2
Static: Random	0.0022	13.3
<b>Thompson Sampling</b>	<b>0.0485</b>	<b>292.7</b>
$\epsilon$ -greedy (0.1)	0.0474	286.1
$\epsilon$ -greedy (0.3)	0.0427	258.0
UCB1	0.0406	245.2

**Thompson Sampling** nearly matches the best static policy (Hybrid: 293.9 vs Thompson: 292.7) without knowing in advance which arm is best. It maintains a Beta( $\alpha, \beta$ ) distribution for each arm — samples a random value from each distribution, then pulls the arm with the highest sample. After a success  $\alpha += 1$ , after failure  $\beta += 1$ . Over time, the distributions narrow around the true reward, naturally reducing exploration.

**$\epsilon$ -greedy (0.1)** permanently wastes 10% of traffic on random exploration — acceptable but suboptimal. At  $\epsilon=0.3$ , 30% random exploration wastes too much traffic and drops total reward by 10%.

**UCB1 failed.** Rewards are small ( $\approx 0.04$ ) while UCB1's uncertainty bonus starts at  $\approx 1.0$  ( $\sqrt{2 \cdot \ln(t)/n}$ ). The bonus dominates arm selection for thousands of rounds, causing severe over-exploration before the average rewards become distinct enough to pick a winner.

## 6. Part 2 Summary Table

Model	NDCG@10	Precision@10	Recall@10	Coverage	Pop.-Bias	Best For
Hybrid Blend ( $\alpha=0.8$ )	0.0657	0.0452	0.0708	48.1%	1058	Cold-start users
BPR-MF	0.0664	0.0466	0.0704	45.9%	1095	Overall ranking
CandGen+Rerank	0.0655	0.0455	0.0701	47.2%	1067	Production architecture
Two-Tower	0.0648	0.0477	0.0666	54.1%	1199	Production retrieval
NeuMF	0.0619	0.0447	0.0622	56.8%	1227	Flexibility
PPR	0.0544	0.0428	0.0497	5.8%	2053	Graph structure
Popularity	0.0490	0.0393	0.0433	5.1%	2086	Baseline
Recency	0.0438	0.0353	0.0410	4.4%	1797	Freshness

## 7. Unified Results: All Models

Rank	Model	NDCG@10	Precision@10	Recall@10	Coverage	Pop. Bias
1	BPR-MF	0.0664	0.0466	0.0704	45.9%	1095
2	Hybrid Blend ( $\alpha=0.8$ )	0.0657	0.0452	0.0708	48.1%	1058
3	CandGen+Rerank	0.0655	0.0455	0.0701	47.2%	1067
4	Two-Tower	0.0648	0.0477	0.0666	54.1%	1199
5	Item-Item CF	0.0625	0.0457	0.0593	19.7%	1681

Rank	Model	NDCG@10	Precision@10	Recall@10	Coverage	Pop. Bias
6	NeuMF	0.0619	0.0447	0.0622	56.8%	1227
7	PPR	0.0544	0.0428	0.0497	5.8%	2053
8	Popularity	0.0490	0.0393	0.0433	5.1%	2086
9	Enhanced CB	0.0460	0.0302	0.0495	40.4%	978
10	Recency	0.0438	0.0353	0.0410	4.4%	1797
11	FunkSVD	0.0265	0.0218	0.0223	16.8%	799
12	ALS	0.0219	0.0173	0.0185	84.5%	470
13	Basic CB	0.0087	0.0064	0.0090	100.2%	237

## 8. Key Insights

- Objective alignment > model complexity.** BPR-MF and FunkSVD are the same architecture. Switching from MSE to pairwise ranking loss: +150% NDCG. This is the single most impactful finding.
- The accuracy–coverage trade-off is real.** Item-Item CF: 0.0625 NDCG, 19.7% coverage. ALS: 0.0219 NDCG, 84.5% coverage. No model wins both.
- Hybrid helps cold-start, hurts warm users.** Blend ( $\alpha=0.8$ ) boosts cold-user NDCG by +10% over BPR. But for warm users, BPR alone is better — content features add noise when collaborative signal is strong.
- Deep learning ≠ automatic improvement.** On MovieLens 1M, linear BPR-MF beats both NeuMF and Two-Tower. Neural networks need more data (or richer features) to justify their extra parameters.
- Offline metrics are insufficient.** Thompson Sampling matched the best static policy without prior knowledge. The exploration it performs creates value invisible to offline evaluation.

## 9. Problems & Struggles

**Objective mismatch (FunkSVD/ALS):** Took real debugging time to understand why MF models had excellent RMSE but terrible NDCG. Not a bug — a fundamental misalignment between training objective and evaluation metric.

**BPR loss sign:** The BPR log-likelihood is negative and increases toward 0 as the model improves. Our training plots initially looked like divergence (loss going "up"). In the PyTorch models (NeuMF, Two-Tower) we negated the loss to get a standard decreasing curve.

**Score normalization in hybrid:** BPR scores are unbounded; CB scores are in [0, 1]. Blending without normalization produces nonsensical results. We used min-max normalization, which works but is sensitive to outliers. The CandGen+Rerank approach sidesteps this by using raw BPR scores with a lower weight (0.6) — it works because within the top-100 BPR candidates, the score range is narrower.

**UCB1 failure:** Specific to our reward scale. Rewards  $\approx 0.04$ , but the UCB1 bonus  $\sqrt{2 \cdot \ln(t)/n}$  starts at  $\approx 1.0$ , causing thousands of rounds of over-exploration before the average rewards become distinct enough to pick a winner. A reward-scaled variant or different confidence parameter would fix this.

**Computational cost:** BPR at 20x sampling = 9.4M samples/epoch. Without Numba JIT, each epoch would take minutes instead of seconds. Neural models required GPU. Item-Item CF similarity is  $O(n\_items^2)$  — several minutes on 3,706 items.

## 10. Deployment Recommendation

### Two-stage architecture:

1. **Retrieval: Two-Tower.** Precompute all 3,706 item representations through the item tower. Index them in FAISS. Per-user inference: one forward pass through the user tower + ANN search → top-100 candidates. Content features in the item tower help with cold-start items.
2. **Reranking: Lightweight Hybrid.** Rerank the 100 candidates using
$$\alpha \cdot \text{retrieval\_score} + (1-\alpha) \cdot \text{business\_signals}$$
. Business signals include: CB genre match, recency boost, diversity constraint, already-seen genre demotion.

### Post-deployment — two stages:

- **Stage 1 — A/B test:** Validate the new pipeline against the Item-Item CF baseline. Primary metric: CTR on top-5 recommendations. Guardrail: Watch Time (to avoid optimizing for clickbait). Requires statistical significance before rollout.
- **Stage 2 — Thompson Sampling:** Once validated, use a bandit to dynamically route traffic between model versions as they evolve. No need to rerun A/B tests for every model update — the bandit converges to the best option automatically.